

1. General context
2. Supervised vs. unsupervised learning
3. Machine learning in Physical Chemistry
4. Python vs. other tools – at the era of generative AIs
5. Typical workflow – including xAI
6. Tutorials / Live demonstrations ← Jupyter notebooks

General context

Artificial Intelligence (AI)

intelligence demonstrated by machines, as opposed to the natural intelligence displayed by humans or animals

Goals

reasoning & (basic) problem solving

knowledge representation

planning: making choices and hierarchy of events

learning (*i.e.* machine learning)

natural language processing

perception of the world from sensors

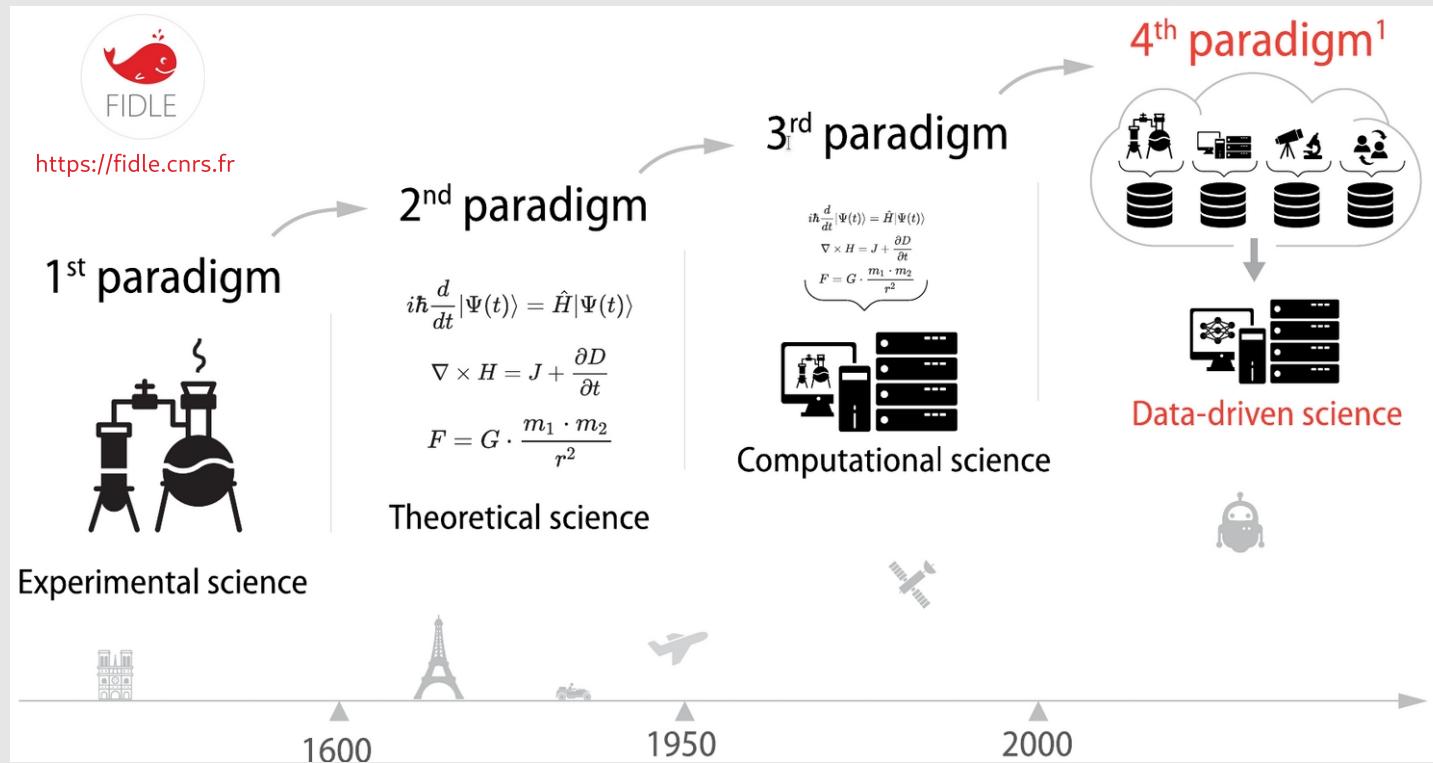
ability to move and manipulate objects

simulating human affects

long-term goal: ability to solve an arbitrary problem

N.B. in some fields "artificial intelligence" means "machine learning with neural networks"

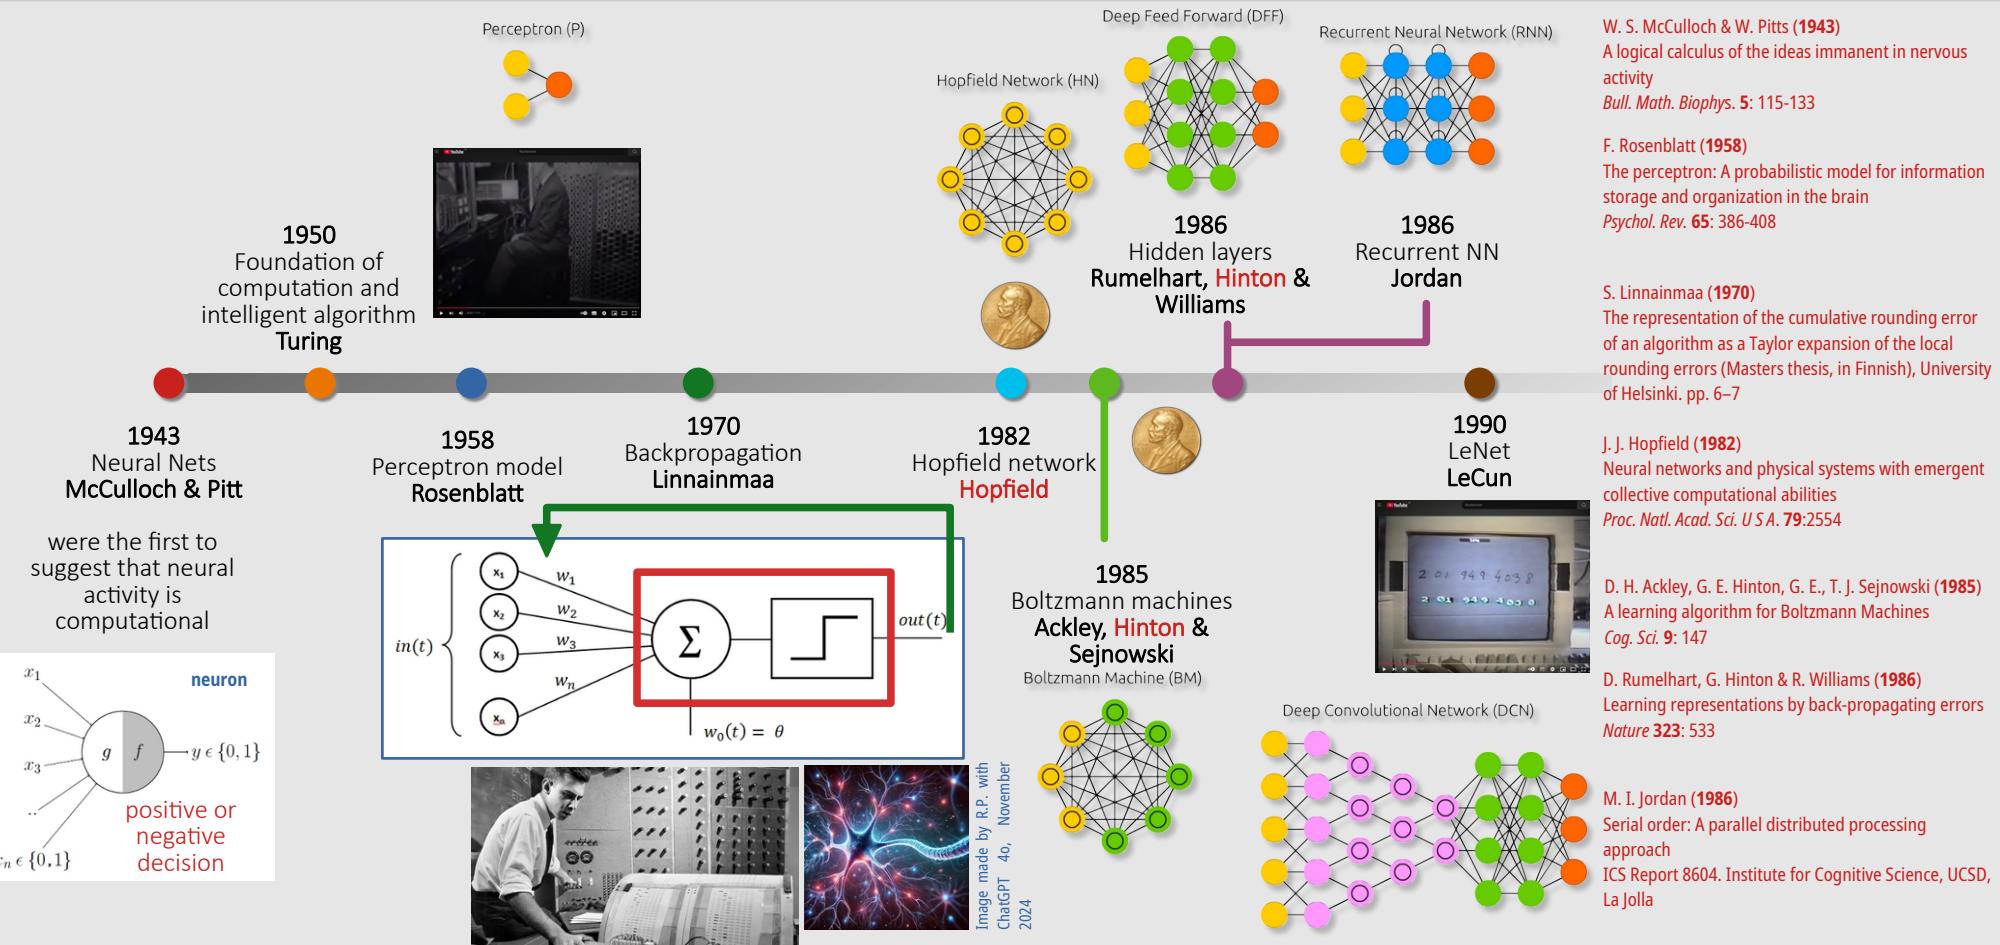
Artificial Intelligence and Scientific Research



Outlook: data science and machine learning

Artificial Intelligence and Artificial Neural Networks (ANN) timeline

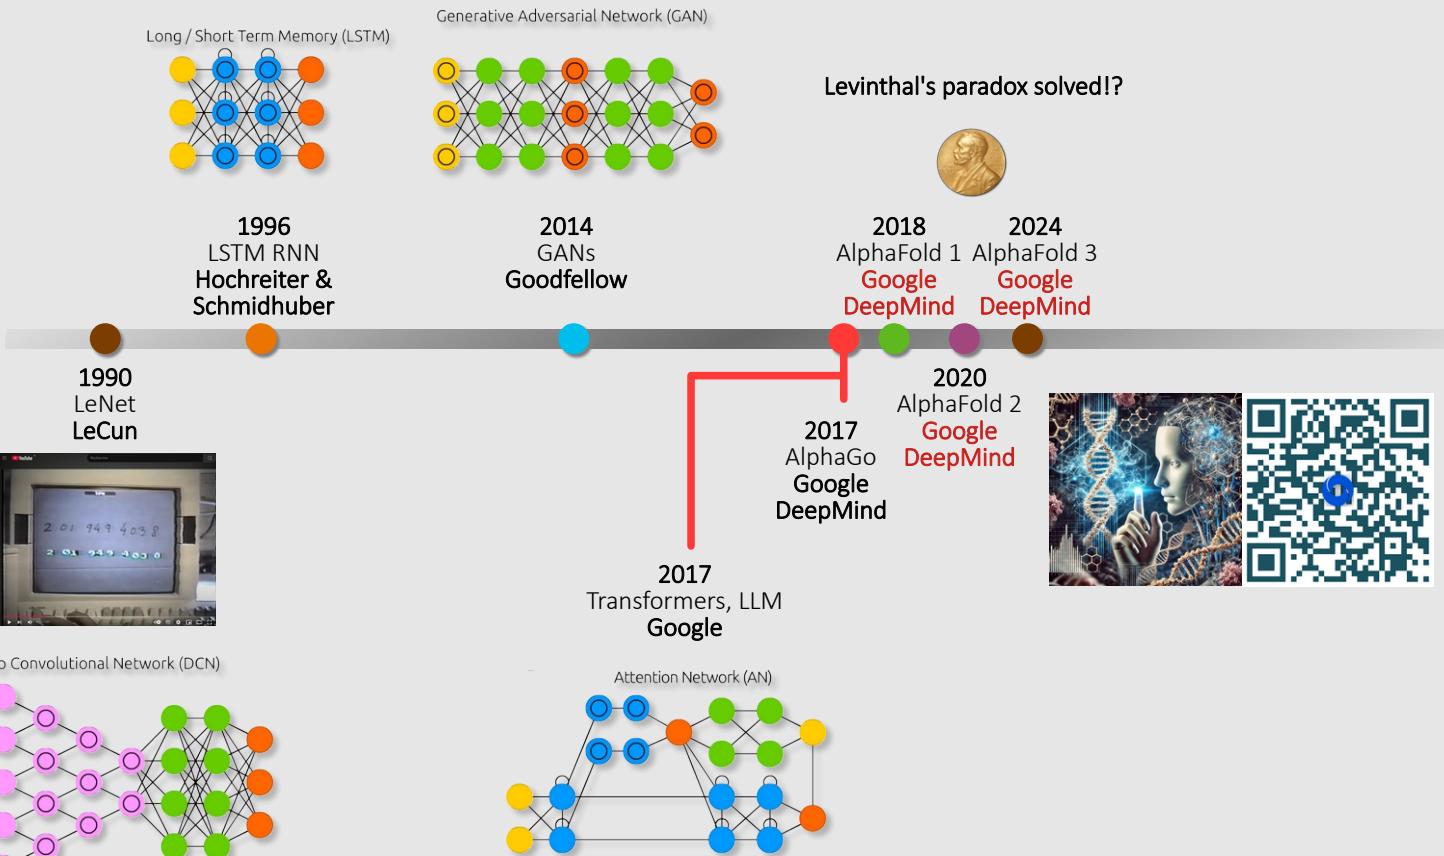
Version: Tuesday, December 2, 2025



Outlook: data science and machine learning

Artificial Intelligence and Artificial Neural Networks (ANN) timeline

Version: Tuesday, December 2, 2025



S. Hochreiter & J. Schmidhuber (1996)
LSTM can solve hard long time lag problems
Advances in Neural Information Processing Systems 9:
Proceedings of The 1996 Conference, MIT Press 1997,
p. 473

I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, Y. Bengio (2014)
Generative Adversarial Nets. Proceedings of the International Conference on Neural Information Processing Systems (NIPS 2014), p. 2672

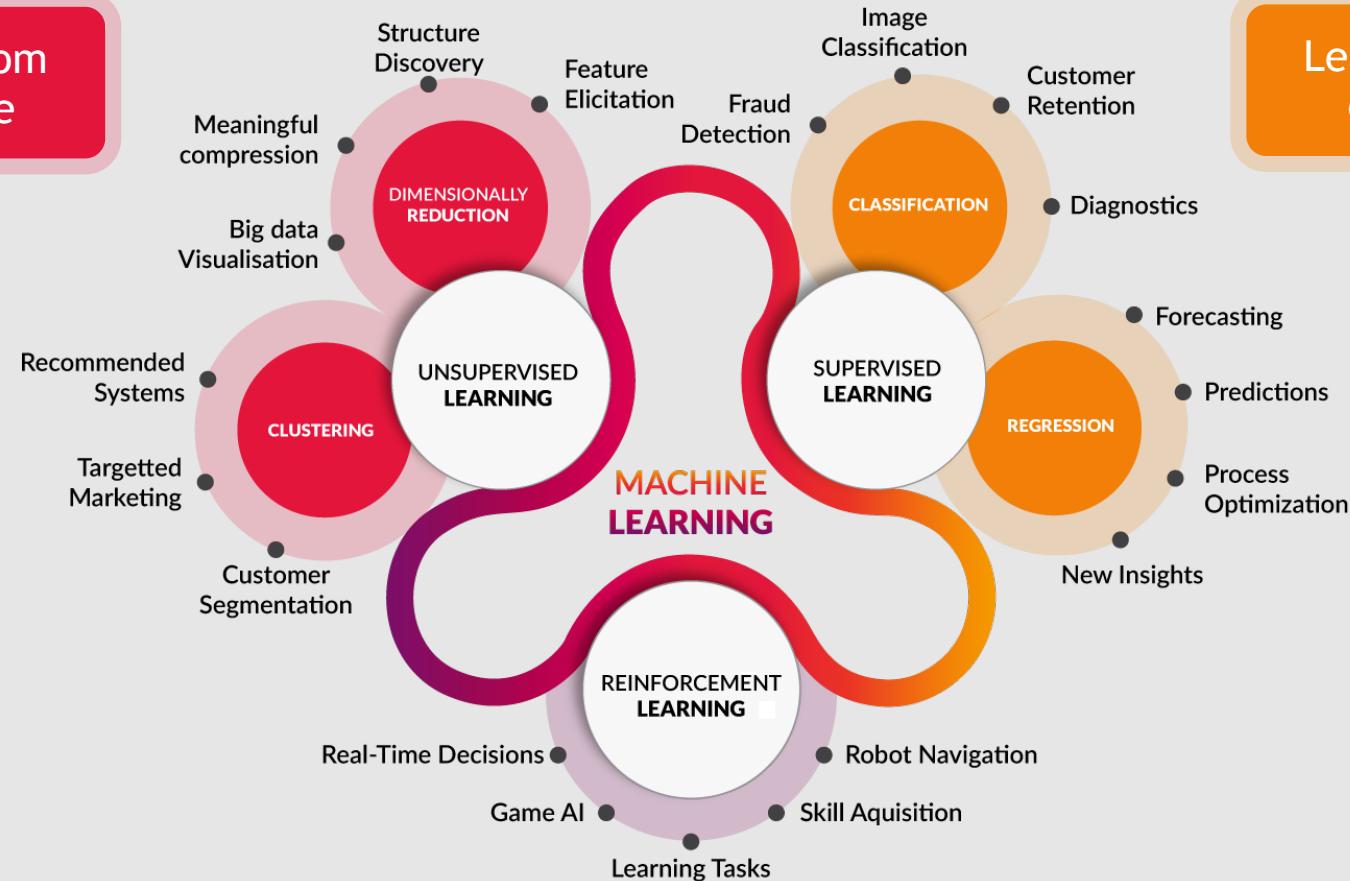
A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, I. Polosukhin (2017)
Attention is All you Need
in Advances in Neural Information Processing Systems. 30. Curran Associates, Inc.

J. Abramson, J. Adler, J. Dunger, R. Evans, T. Green, A. Pritzel, O. Ronneberger, L. Willmore, A. J. Ballard, J. Bambrick, S. W. Bodenstein, D. A. Evans, C.-C. Hung, M. O'Neill, D. Reiman, K. Tunyasuvunakool, Z. Wu, A. Žemgulytė, E. Arvaniti, C. Beattie, O. Bertoli, A. Bridgland, A. Cherepanov, M. Congreve, A. I. Cowen-Rivers, A. Cowie, M. Figurnov, F. B. Fuchs, H. Gladman, R. Jain, Y. A. Khan, C. M. R. Low, K. Perlin, A. Potapenko, P. Savy, S. Singh, A. Stecula, A. Thillaisundaram, C. Tong, S. Yakneen, E. D. Zhong, M. Zielinski, A. Žídek, V. Bapst, P. Kohli, M. Jaderberg, D. Hassabis & J. M. Jumper (2024)
Accurate structure prediction of biomolecular interactions with AlphaFold 3
Nature 630: 493

Machine Learning

Machine learning

Learning from data alone



Learning from examples

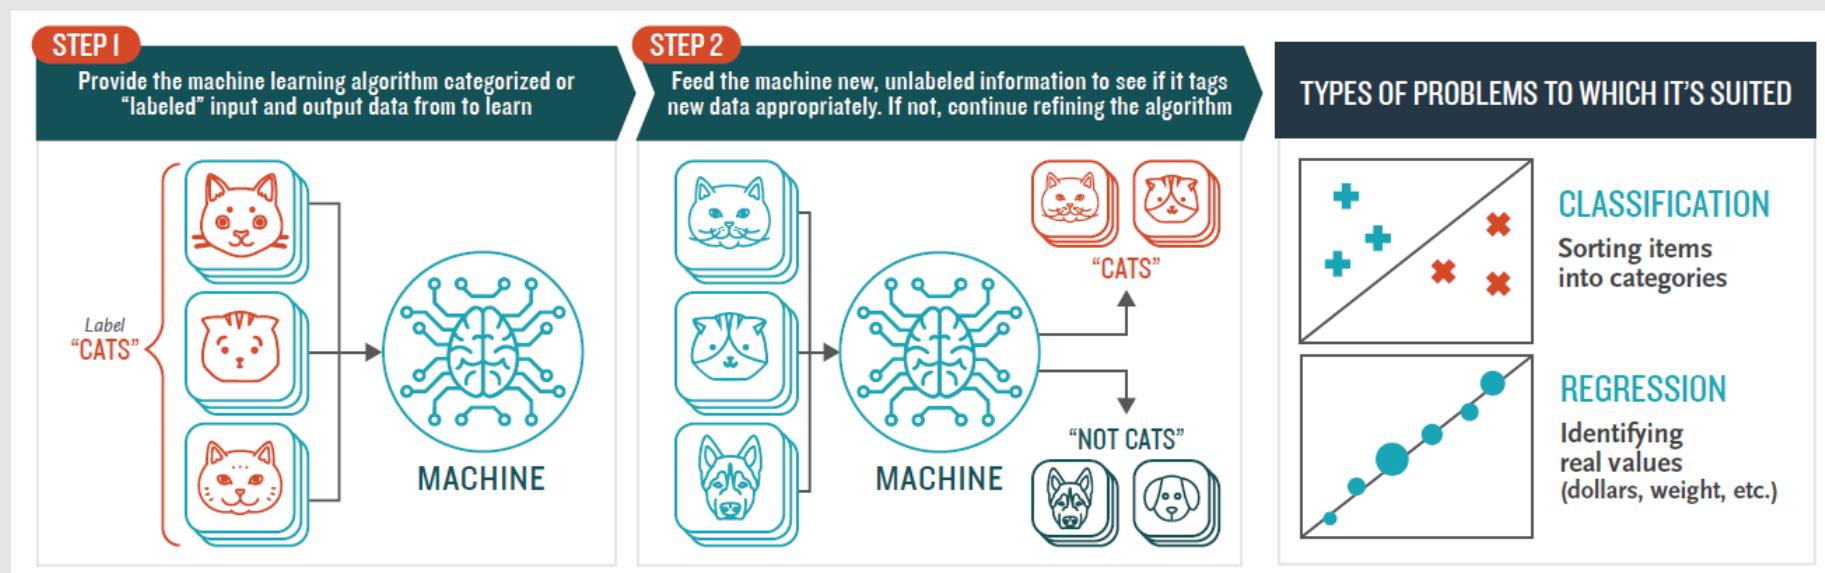
Supervised learning

Data are provided along with the desired output (*i.e.* labelled data)

Example of cats detection:

- collect thousands of images of cats
- draw a bounding box around each cat
- feed the entire dataset to the machine so it can learn all by itself

Learning from examples



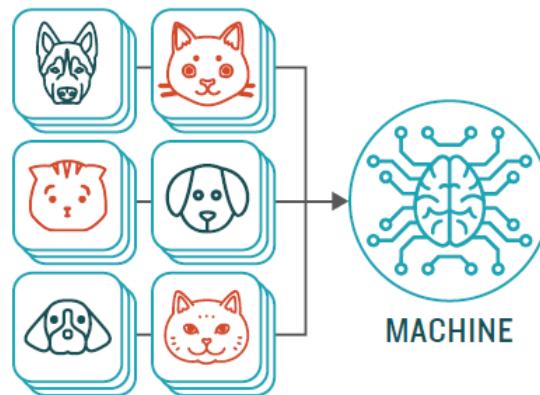
Unsupervised learning

Learning from data alone

- Just provide data
- Let the machine find out (or cluster) the patterns in the dataset

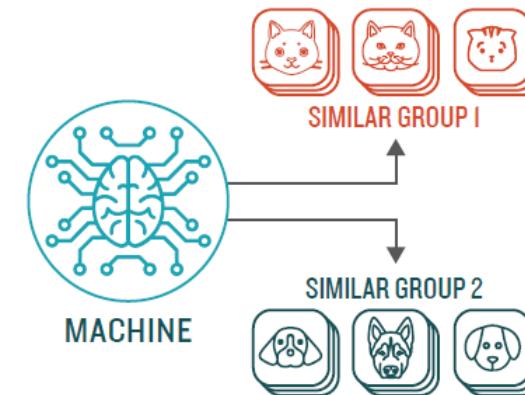
STEP 1

Provide the machine learning algorithm uncategorized, unlabeled input data to see what patterns it finds



STEP 2

Observe and learn from the patterns the machine identifies

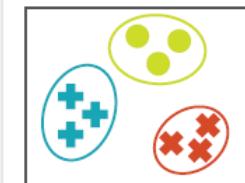


TYPES OF PROBLEMS TO WHICH IT'S SUITED

CLUSTERING

Identifying similarities in groups

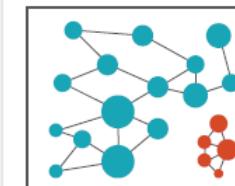
For Example: Are there patterns in the data to indicate certain patients will respond better to this treatment than others?



ANOMALY DETECTION

Identifying abnormalities in data

For Example: Is a hacker intruding in our network?



Machine Learning in Physical Chemistry

Transformative role of AI in sciences of matter research



Image made by R.P. with ChatGPT 4o, November 2024

Acceleration of Material Discovery

- Property prediction prior to laboratory synthesis
- Exploration of a broad parameter space
- Inverse design of materials
- Catalyst design

Understanding mechanisms and multi-factorial processes through explainable AI (XAI)

Advanced Modeling

- Predictions comparable to *ab initio* calculations
- Simulations of physical and chemical processes at challenging scales

Design of Experiments

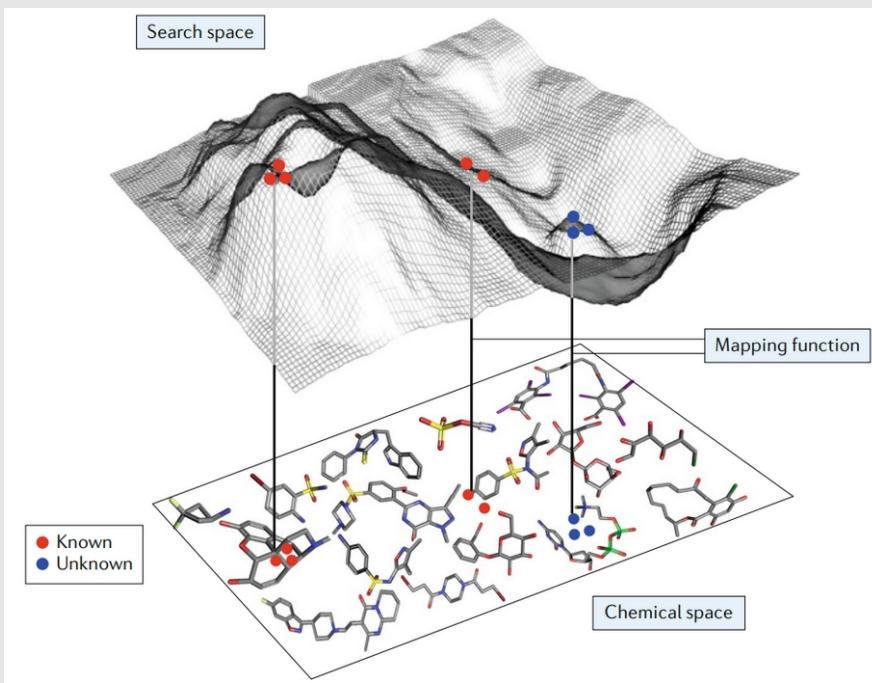
- Automated processing with robotic systems guided by AI and real-time analysis of experimental data
- Optimizing experimental setups through Design of Experiments (DoE)

Data Exploration

- Detecting trends and anomalies using unsupervised methods
- Derive insights from diverse datasets like crystal structures or spectroscopy

The chemical space is unimaginably vast

~ 10^{60} drug-like molecules



Experimental and quantum chemistry methods:
accurate but slow or costly

Most of the chemical space remains unexplored

The plan

- to exploit available data (experimental or DFT, small or large)
- to train data-driven models using ML
- and use explainability tools to open the black box of predictions



Development of AI tools with python

How to develop home-made ML tools?

Version: Tuesday, December 2, 2025



The ML galaxy within python™ is an exceptional ecosystem

Python is a high-level, interpreted, object-oriented, general-purpose programming language



<https://scikit-learn.org/>



<https://pytorch.org/>



<https://keras.io/>



NumPy



seaborn



pyNMB



<https://www.tensorflow.org/>

Using Python as an everyday tool for scientific calculation and computation (it can even replace excel... by far)

- basic knowledge of programming languages (variables, arrays, loops, conditional tests...)
- enthusiastic and vast community



→ cheatsheets & cut/paste



The image displays four separate screenshots of Python cheat sheets from Dataquest, arranged in a 2x2 grid. Each screenshot is a dense, multi-colored reference guide containing numerous code snippets, tables, and explanatory text. The top-left sheet is titled 'Python 3 Cheat Sheet' and covers basic operations like arithmetic, comparison, and logical operators. The top-right sheet is titled 'Python For Data Science Cheat Sheet - Python Basics' and includes sections on variables and data types, array operations, and data structures. The bottom-left sheet is titled 'Data Science Cheat Sheet - Beginner' and focuses on data manipulation with pandas and NumPy. The bottom-right sheet is titled 'Data Science Cheat Sheet - Intermediate' and covers more advanced topics like machine learning and data visualization with Matplotlib and Seaborn.

Uneasy?

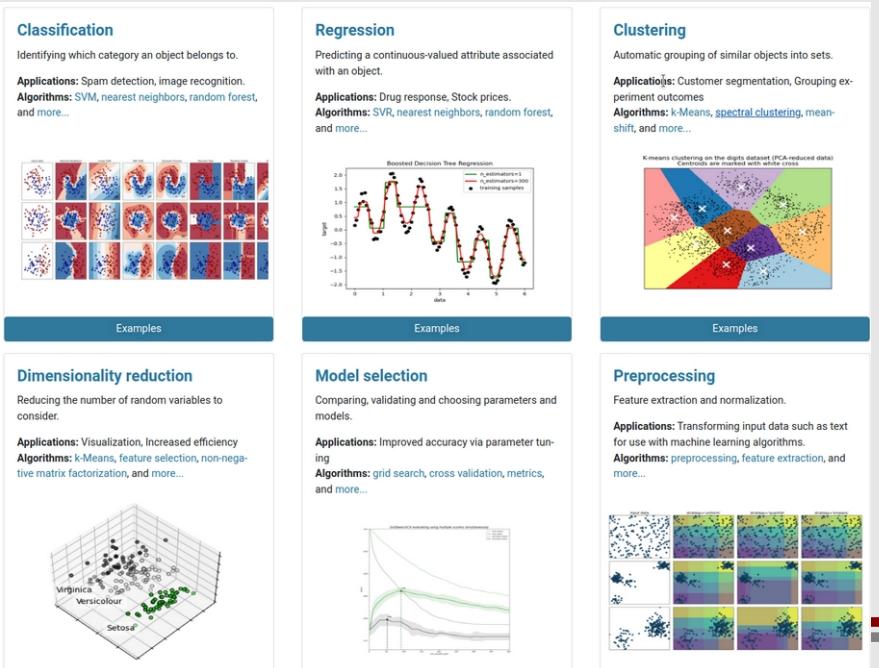
- yes and no
- important initial investment
- worth the effort



Machine learning in Python with scikit-learn

Simple and efficient tools for predictive data analysis
Accessible to everybody, and reusable in various contexts
Built on NumPy, SciPy, and matplotlib

INRIA took leadership of the project and made the first public release on February 2010
3-clause BSD License (permissive free software license, compatible with the GNU GPL)



Classification
Identifying which category an object belongs to.
Applications: Spam detection, image recognition.
Algorithms: SVM, nearest neighbors, random forest, and more...

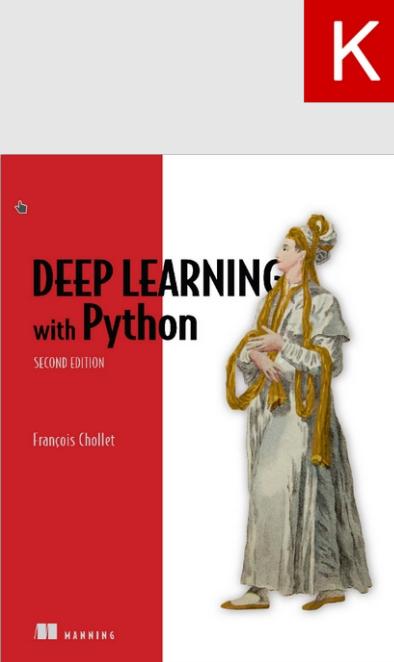
Regression
Predicting a continuous-valued attribute associated with an object.
Applications: Drug response, Stock prices.
Algorithms: SVR, nearest neighbors, random forest, and more...

Clustering
Automatic grouping of similar objects into sets.
Applications: Customer segmentation, Grouping experiment outcomes
Algorithms: k-Means, spectral clustering, mean-shift, and more...

Dimensionality reduction
Reducing the number of random variables to consider.
Applications: Visualization, Increased efficiency
Algorithms: k-Means, feature selection, non-negative matrix factorization, and more...

Model selection
Comparing, validating and choosing parameters and models.
Applications: Improved accuracy via parameter tuning
Algorithms: grid search, cross validation, metrics, and more...

Preprocessing
Feature extraction and normalization.
Applications: Transforming input data such as text for use with machine learning algorithms.
Algorithms: preprocessing, feature extraction, and more...



High Level Deep Learning Application Programming Interface (API)

By François Chollet (Google)
Part on TensorFlow since 2017

MIT license (permissive free software license)

how to start?

```
import numpy as np
import tensorflow as tf
from tensorflow import keras
```



TensorFlow

<https://www.tensorflow.org/>

Google Brain's second-generation system

Supported by Google
Low level API

Apache license (yet another permissive free software license)



<https://pytorch.org/>

From Torch library

Supported by Facebook

BSD licence

(permissive free software license)

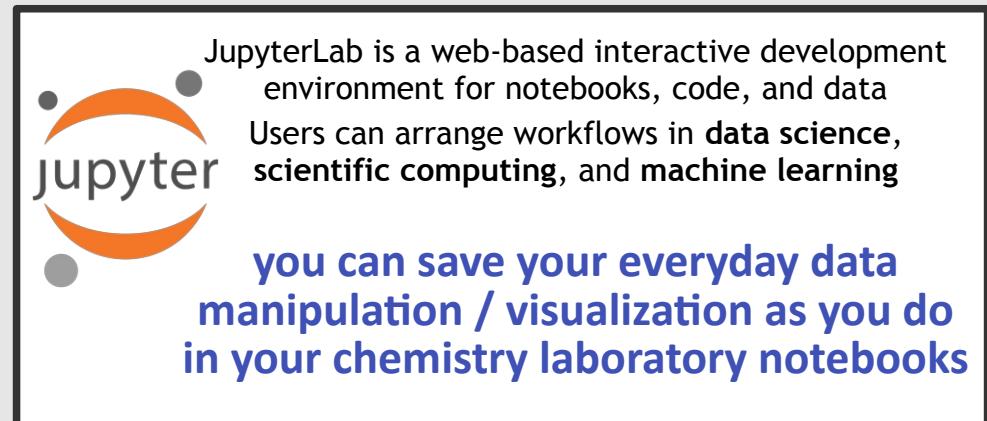
**An open source machine learning framework that accelerates the path
from research prototyping to production deployment**

Widely used in the field of AI research



The screenshot shows the Anaconda website homepage. At the top, there's a navigation bar with links for Products, Pricing, Solutions, Resources, Partners, Blog, and Company. A "Contact Sales" button is also present. Below the navigation, it says "Individual Edition is now ANACONDA DISTRIBUTION". The main heading is "ANACONDA DISTRIBUTION" in large green letters. Below that, it says "The world's most popular open-source Python distribution platform". To the right, there's a callout box for "Anaconda Distribution" with a "Download" button, "For Linux" note, and "Python 3.9 + 64-Bit (x86) Installer • 659 MB" link. It also shows "Get Additional Installers" with icons for Windows, Mac, and Linux.

Anaconda is a distribution of the Python and R programming languages for scientific computing (data science, machine learning applications, large-scale data processing, predictive analytics, etc.), that aims to simplify package management and deployment



JupyterLab is a web-based interactive development environment for notebooks, code, and data. Users can arrange workflows in **data science**, **scientific computing**, and **machine learning**.

you can save your everyday data manipulation / visualization as you do in your chemistry laboratory notebooks

Not familiar with python... yet?

"Python in the Physical Chemistry Lab (PPCL)" in a nutshell

This Python computer lab assumes a very basic knowledge of a programming language and algorithm development.

To run the content of a Python cell: click on a cell to select it. Then press SHIFT+ENTER on your keyboard, or press the play button in the top left corner of this window



If you click on a text cell by accident, you will see the so-called markdown coding of this cell (it is closely related to the HTML language). The corresponding formatted text/images/tables will be rendered by running the cell (SHIFT-ENTER or play button).

Ready? Put down your mobile phone 📱, please, and let's enter into the Python realm. 🎉🎉

PPCL.ipynb

Simple calculations

Basic mathematical operations

```
# Every line that starts with a # character is a comment  
  
# addition  
3 + 2 # it is also possible to add a comment after a command  
  
#This is a new cell. It is possible to define several operations or commands in a cell  
# multiplication  
3*2  
  
#division  
7/2
```

github repository

<https://github.com/rpoteau/pyPhysChem>



README

  Université de Toulouse  [pyPhysChem] 

Ce dépôt GitHub propose une collection de notebooks Jupyter conçus pour intégrer la programmation en Python dans l'enseignement de la chimie physique. Ces notebooks fournissent des exemples commentés et illustrés, couvrant des sujets tels que les dérivées et les intégrales, l'atome d'hydrogène et les représentations moléculaires. Ce dépôt inclut également des ressources pour des applications d'apprentissage automatique en chimie, comme les réseaux de neurones artificiels et les autoencoder. Pour utiliser ces outils, les utilisateurs sont invités à installer Jupyter ainsi qu'une distribution Python, Anaconda étant recommandée. Des instructions détaillées pour cloner le dépôt et exécuter les notebooks sont disponibles dans ce fichier README.md

This GitHub repository offers a collection of Jupyter Notebooks designed to integrate Python programming into physical chemistry education. These notebooks provide commented and illustrated examples, covering topics such as derivatives and integrals, the hydrogen atom, and molecular representations. The repository also includes resources for machine learning applications in chemistry, like artificial neural networks and autoencoders. To utilize these materials, users are advised to install Jupyter and a Python distribution, with Anaconda being a recommended option. Detailed instructions for cloning the repository and running the notebooks are provided in the present README.md document

Table des Matières

- Document principal et pré-requis
- Installation et activation d'une distribution python
 - Introduction
 - Installation de miniconda
 - Activation d'un environnement conda
 - Installation des bibliothèques Python et des outils additionnels nécessaire
- Clonage du dépôt (repository) pyPhysChem et installation des bibliothèques Python nécessaires
- Utiliser ces notebooks à l'aide de JupyterLab
- Liste des changements
- Comment citer ce travail ?

Table of Contents

- Main document and prerequisites
- Installation and activation of a Python distribution
 - Introduction
 - Installing miniconda

gAIs provide assistance to the development of simple or even advanced python codes

Opportunities / Useful Applications

Rapid prototyping of scientific workflows

Data parsing and mining, visualization, automation, small ML pipelines

Assistance for complex coding tasks

DFT parsing, spectroscopy analysis, HPC workflows, error-handling...

Automatic documentation & code explanations

Docstrings, comments, restructuring legacy scripts

Debugging support

Identifying logical errors, improving robustness, optimizing performance

Bridging multi-disciplinary knowledge

Linking chemistry concepts with algorithmic implementations

gAIs can accelerate code development, but the scientist remains responsible for validation, interpretation, and critical checking and thinking

Risks / Limitations

Hallucinations and subtle errors

Syntactically correct but chemically or physically wrong code

Overtrust in outputs

Risk of using unverified code in production or scientific publications

Data leakage / confidentiality concerns

(if code or sensitive data are pasted in public LLMs)

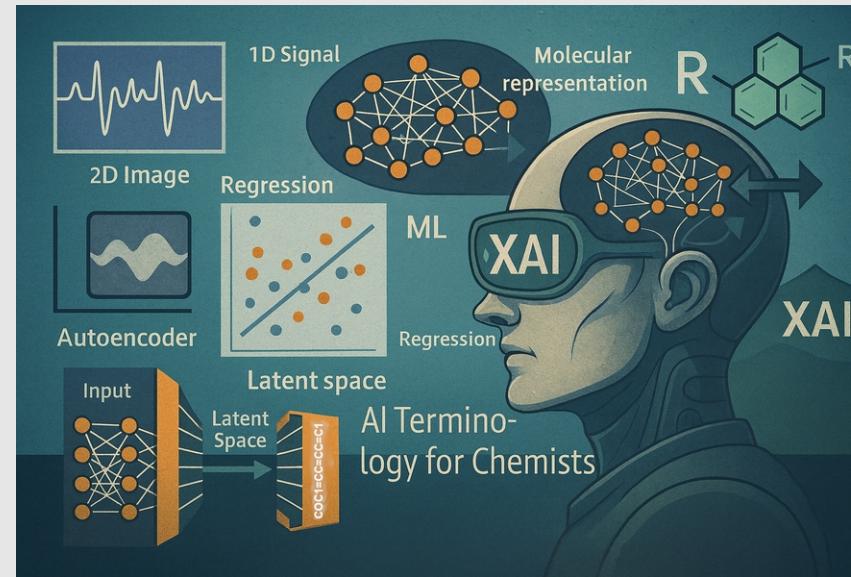
Environmental impact

Training and inference of large models have a non-negligible carbon footprint

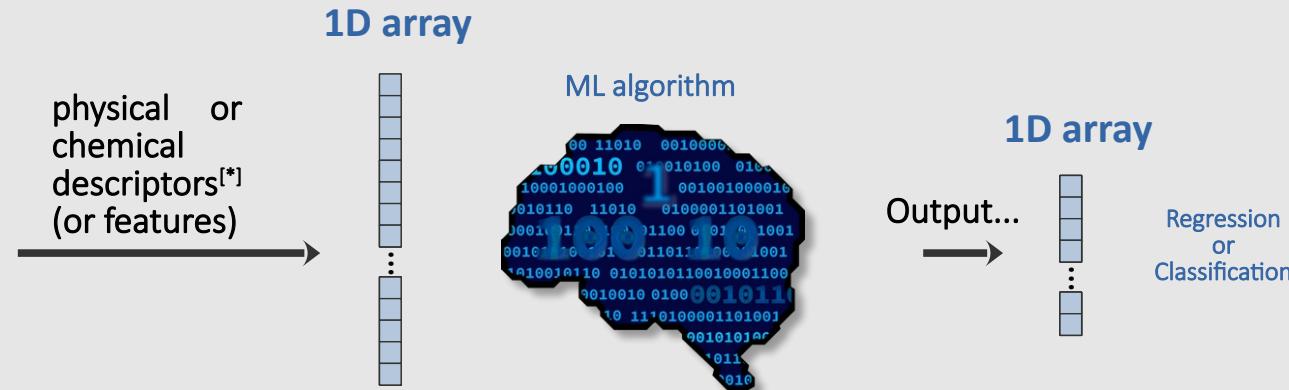
Loss of skill development

If over-used as a substitute for critical thinking instead of a tool

Typical workflow – with a bit of xAI



General context



[*] electron-counting rules, local aromaticity, HOMO-LUMO gaps, d-band centers, Fermi energy, global or local structural parameters, atomic charges, electron density, BDE, experimental spectrum, kinetic curve, thermal conductivity, heat capacity, resistivity, ductility,....

Strategy



Collect data

Which data ?



How much?



Explainable AI tools
to see through the
black box models

What did the
model learn?



Analyze and format
data



Hidden
patterns?

Validate model
performance
on new data



Train and optimize
models

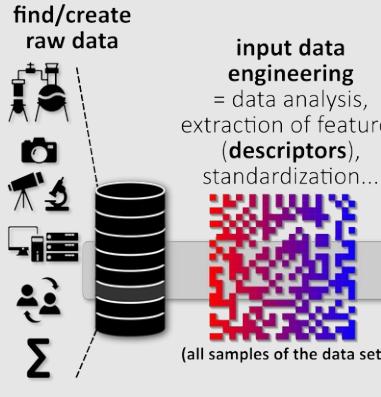
Which
algorithm?



Simplified – but recommended – workflow

Data analysis & feature engineering

Version: Tuesday, December 2, 2025

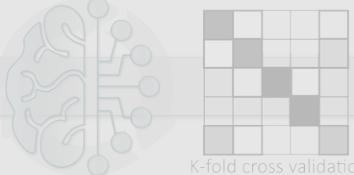


splitting of the dataset (holdout method)
20% test set
80% training set

choice of an ML model and of a metric + hyperparameters tuning

possible need to design a new model:
new data or optimization of hyperparameters or different algorithm)

training, evaluation and optimization of the hyperparameters of the model + validation step



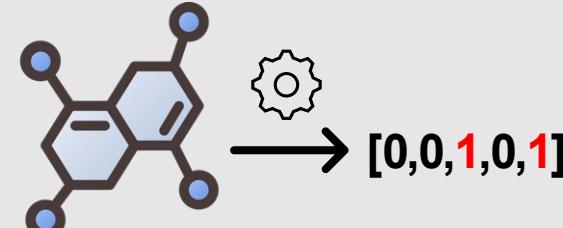
application of the best ML model to new data



xAI = model explainability
Shap

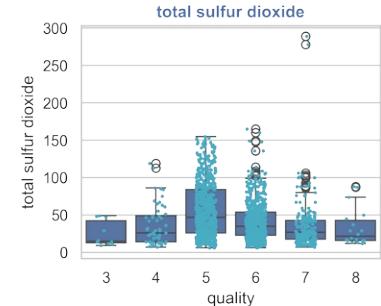


format and transform the raw data into a machine-readable form



Molecular case: conversion of molecular structures into descriptors that numerically represent their features, e.g. vectors

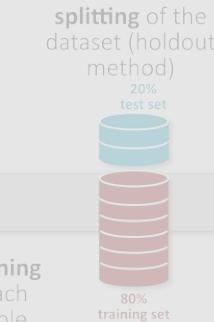
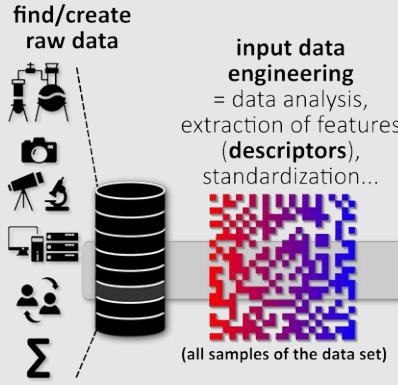
Exploratory Data Analysis Understanding the Dataset Data Quality Assessment



Simplified – but recommended – workflow

Data analysis & feature engineering

Version: Tuesday, December 2, 2025



choice of an ML model and of a metric + hyperparameters tuning

possible need to design a new model:
new data or optimization of hyperparameters or different algorithm)

training, evaluation and optimization of the hyperparameters of the model + validation step



K-fold cross validation

application of the best ML model to new data



xAI = model explainability
Shap

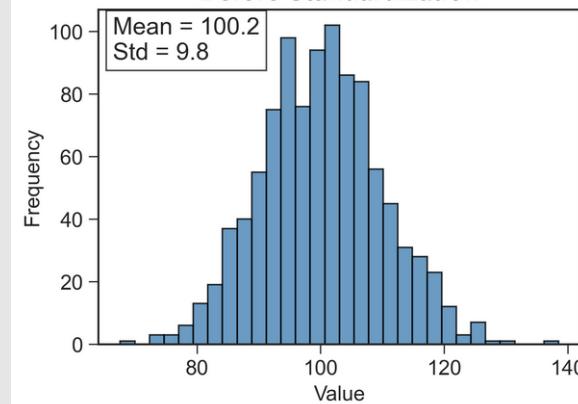


format and transform the raw data into a machine-readable form

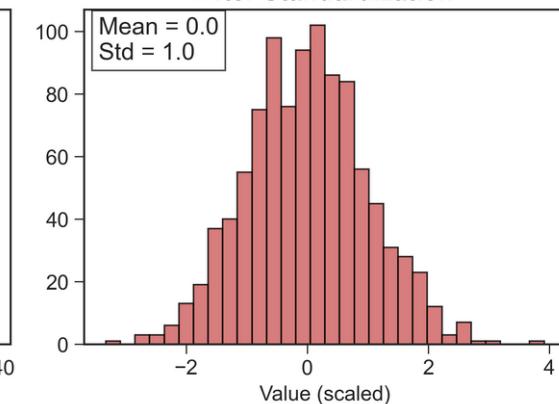


data standardization

Before Standardization

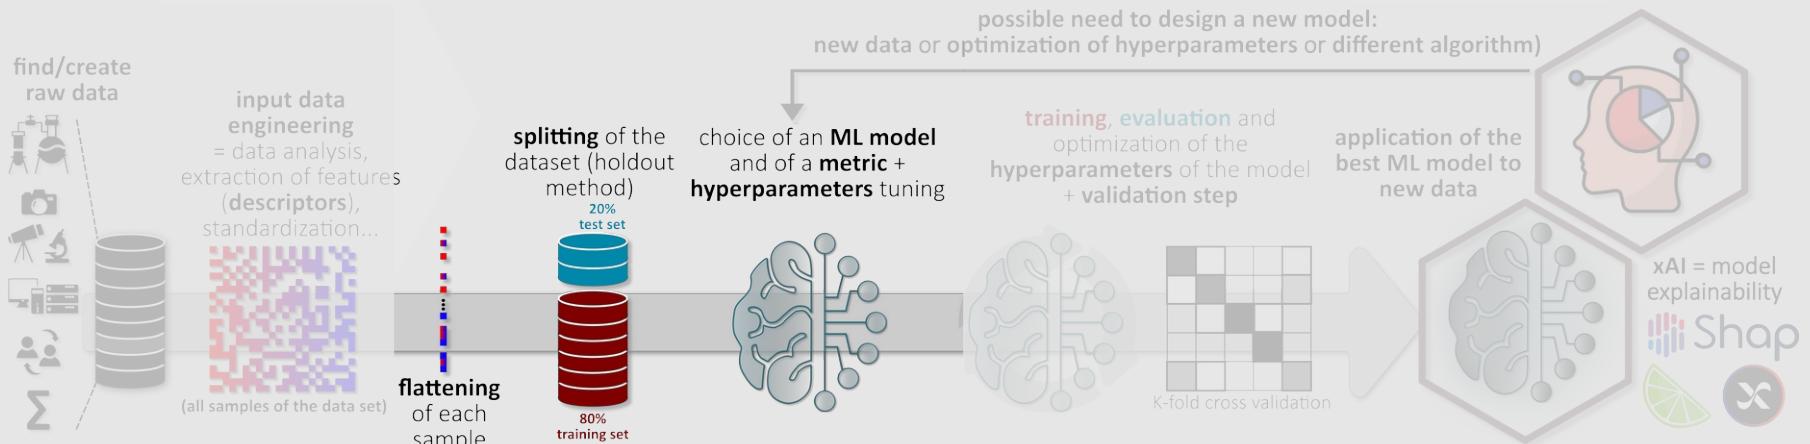


After Standardization



Simplified – but recommended – workflow

Choice of an algorithm & of its hyperparameters



Supervised learning

we want the model to learn the relationship between input features/descriptors (X) and outputs - target (Y)

We split the data:

- 80 % for training → to teach the model
- 20 % for test → to check how well it learned on new data

Hyperparameters? e.g. number of neurons, number of neuron layers...

Usual metrics
(=loss function)

$$\mathcal{L}_{\text{MAE}} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

$$\mathcal{L}_{\text{MSE}} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

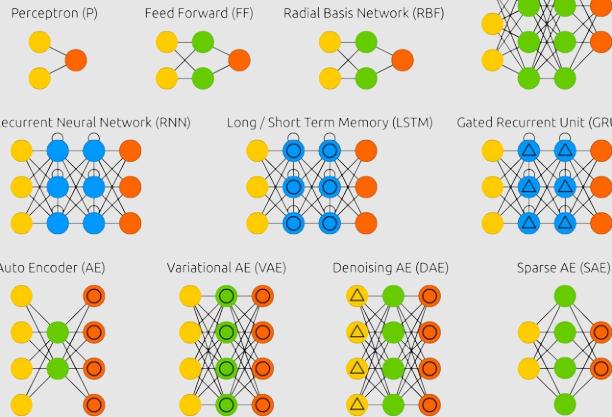
Simplified – but recommended – workflow

Choice of an algorithm: zoo of neural networks (NN)

- Input Cell
- Backfed Input Cell
- △ Noisy Input Cell
- Hidden Cell
- Probabilistic Hidden Cell
- △ Spiking Hidden Cell
- Capsule Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- △ Gated Memory Cell
- Kernel
- Convolution or Pool

A mostly complete chart of Neural Networks

©2019 Fjodor van Veen & Stefan Leijnen asimovinstitute.org



Simplified – but recommended – workflow

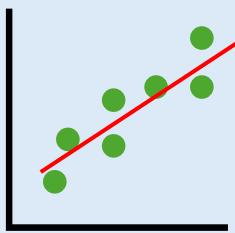
Choice of an algorithm: neuron-free algorithms

Linear methods

Fit a straight line through the data

Simple, fast, interpretable

Only for linear relationships



Examples:

- Linear regression
- Ridge regression

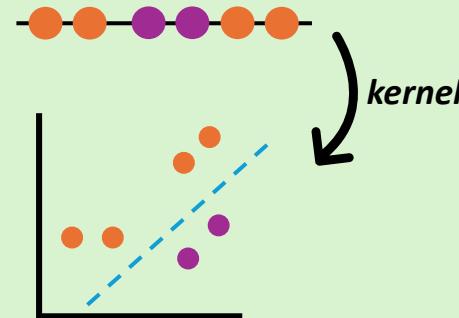


Kernel-based methods

Use a mathematical trick to capture non-linear patterns

Flexible, good for small datasets

Computationally expensive for large data



Examples:

- Support Vector Machines
- Support Vector Regressors

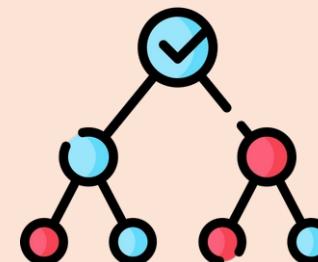


Tree-based methods

Make rule-based “if–then” decisions

Power + interpretability, good for small datasets

Cannot extrapolate, instability



Examples:

- Decision Trees
- Random Forests

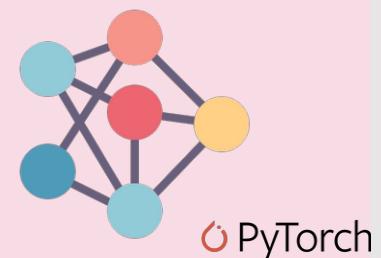


Neural networks

Stack layers of “neurons” to learn complex patterns

Very powerful for large datasets

Need a lot of data, overfitting risk



Examples:

- Deep Neural Networks
- Autoencoders...

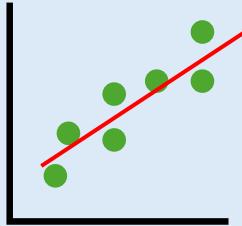


Simplified – but recommended – workflow

Version: Tuesday, December 2, 2025

Linear methods

Fit a straight line through the data



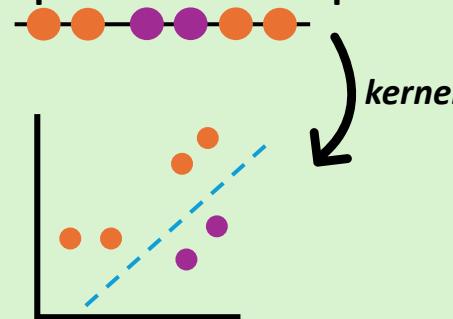
Examples:

- Linear regression
- Ridge regression



Kernel-based methods

Use a mathematical trick to capture non-linear patterns



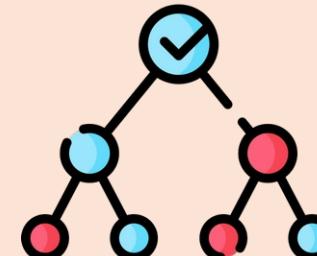
Examples:

- Support Vector Machines
- Support Vector Regressors



Tree-based methods

Make rule-based “if–then” decisions



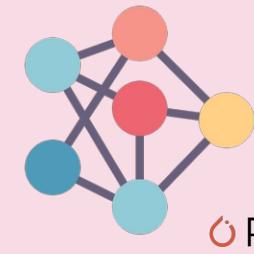
Examples:

- Decision Tree
- Random Forest



Neural networks

Stack layers of “neurons” to learn complex patterns



Examples:

- Deep Neural Networks
- Autoencoders



Bagging (Bootstrap Aggregating)

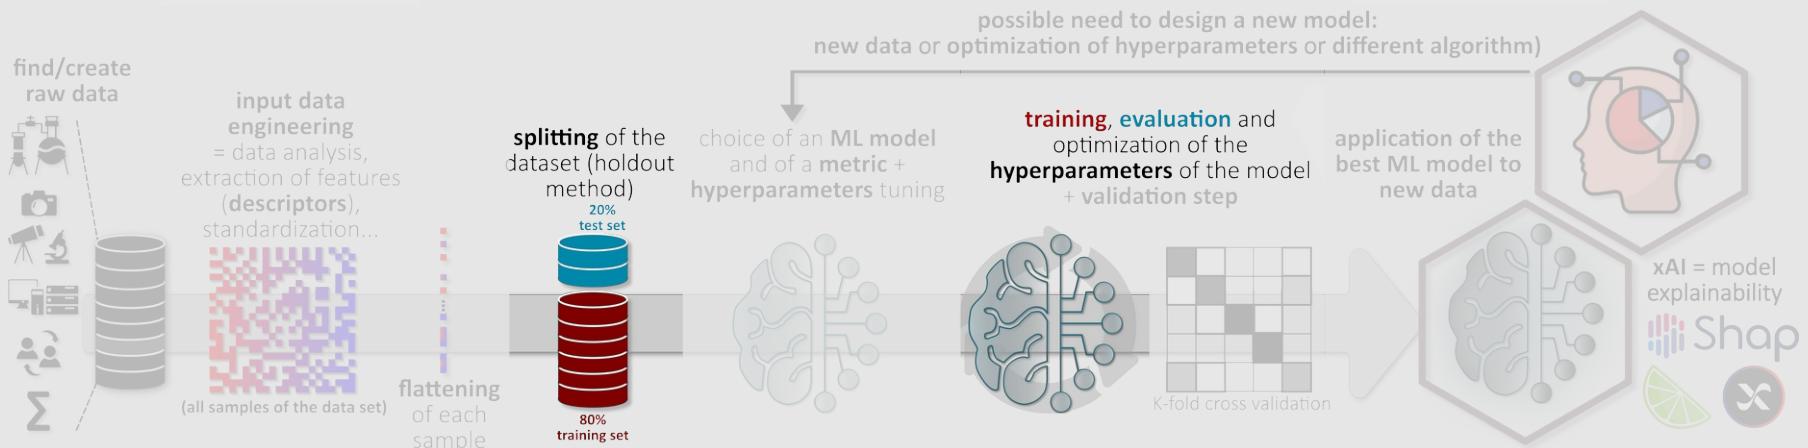
Train several models on random subsets of the data. Then average their predictions.

Example: SVR-bagging = multiple SVRs trained on different samples, predictions averaged.



Simplified – but recommended – workflow

Model training – Supervised learning



We want the model to learn the relationship between input features/descriptors (X) and outputs - target (Y)

The model makes predictions \hat{Y}

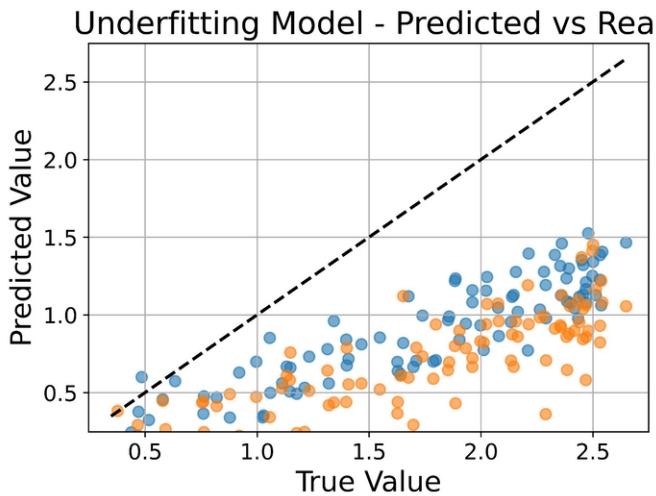
Learning process: minimizing the error between actual values (Y) and predicted values (\hat{Y}) by adjusting the model's internal parameters

It can be iterative (NN case) or analytical (linear least square fit case)

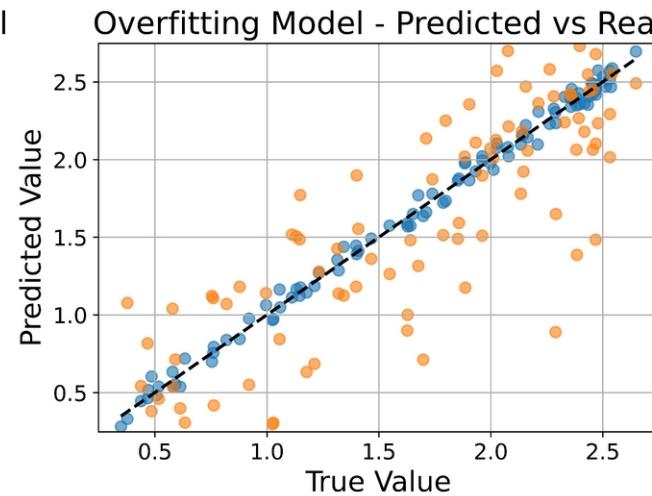
Simplified – but recommended – workflow

Model evaluation – at the end of the training process. Regression case

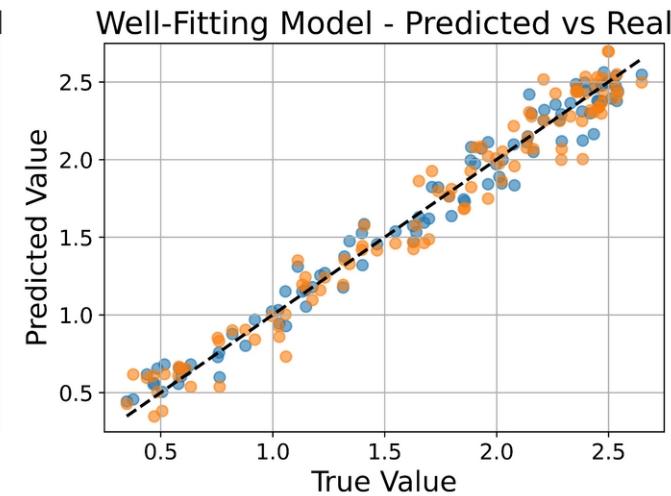
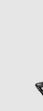
The model is too simple, fails to capture patterns in the training data → poor on both training and test sets



The model is too complex, memorizes the training data → good on training but poor on new data



The model captures the main trends in the data → good performance and generalization



Performance can possibly be improved by changing the hyperparameters of the model

Simplified – but recommended – workflow

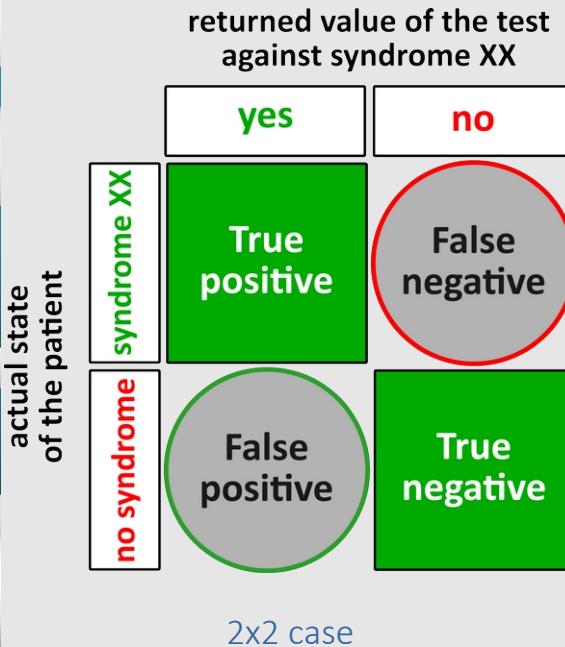
Model evaluation – at the end of the training process. Classification case

consists in plotting confusion matrices (CMs)

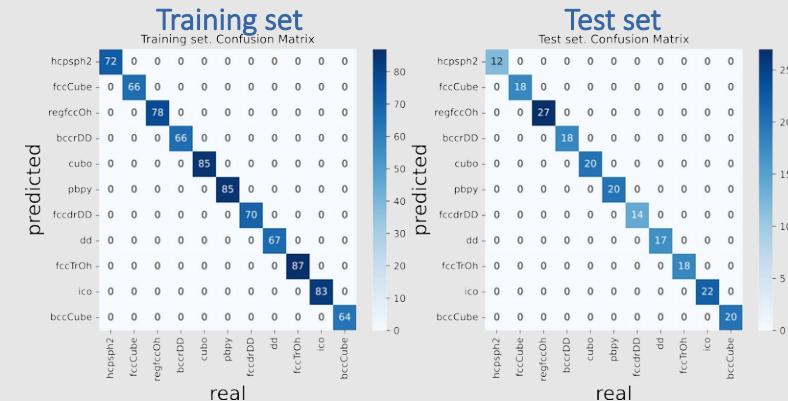
A confusion matrix is a 2×2 table (for binary classification) comparing the model's predictions to the true labels

Same logic as a biological/medical diagnostic test for a syndrome

- A **sensitive** test minimizes false negatives → good for dangerous syndromes
- A **specific** test minimizes false positives → good when false alarms are costly
- ML models behave exactly like diagnostic tests: threshold tuning shifts balance between sensitivity and specificity



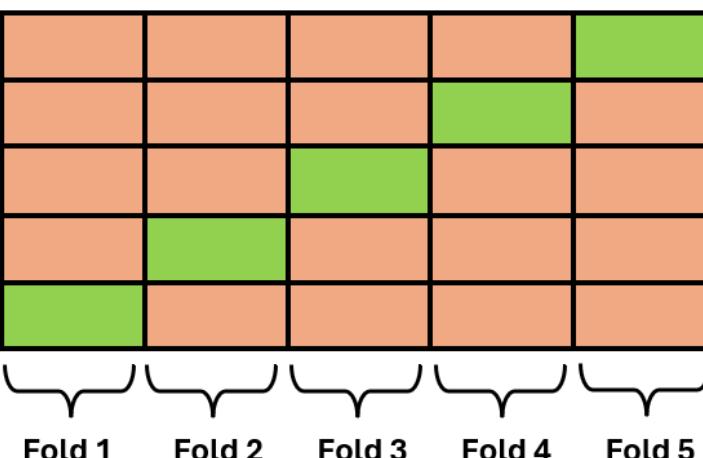
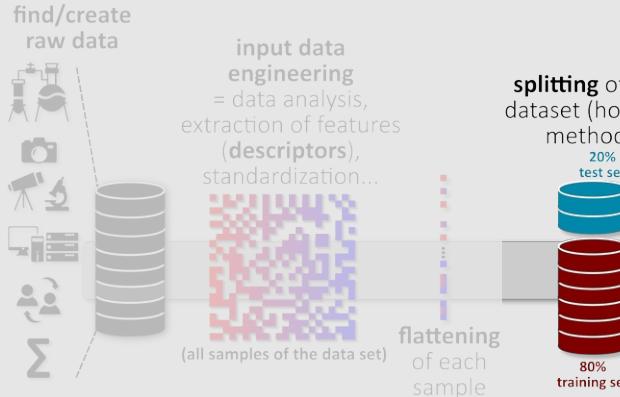
CM generalizes naturally to multi-class classification



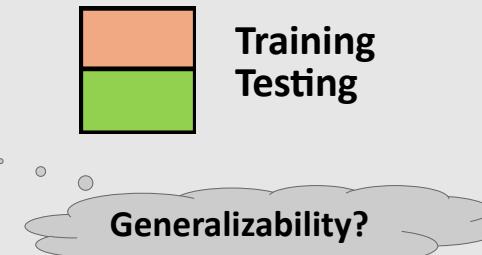
Simplified – but recommended – workflow

Version: Tuesday, December 2, 2025

Model validation



Cross-Validation



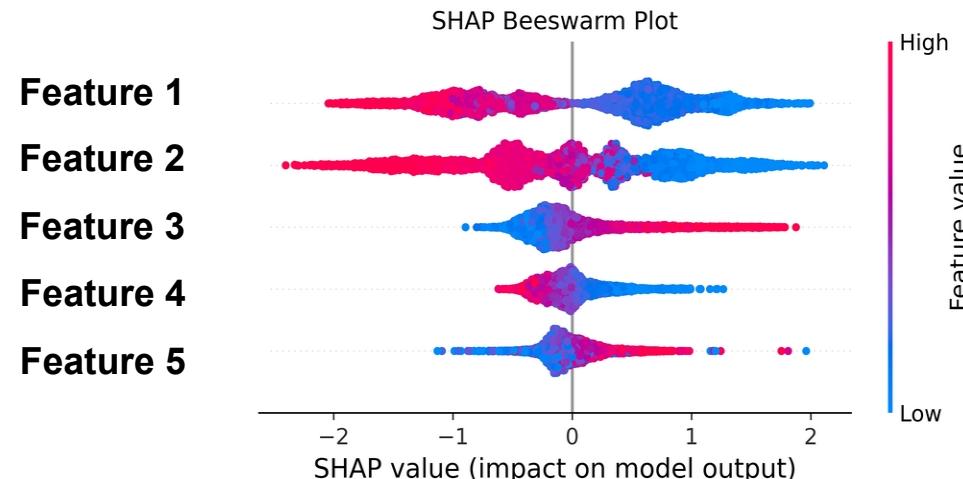
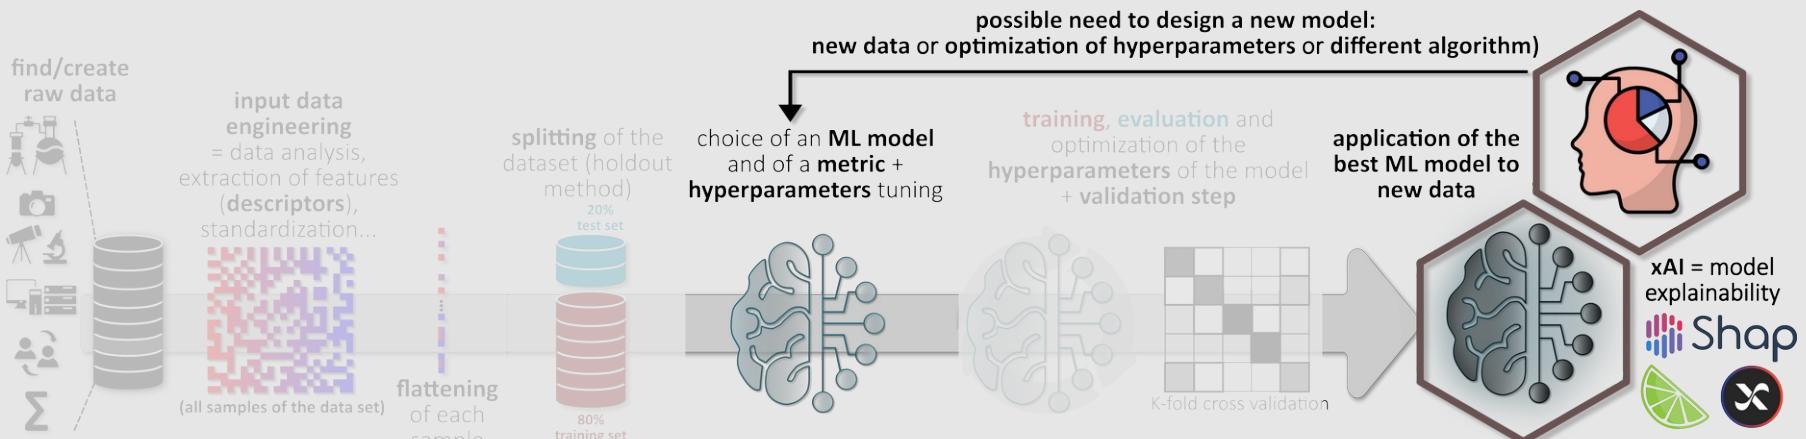
Extrapolation

New “challenging” sets of data



Simplified – but recommended – workflow

Model explainability = eXplainable AI (xAI)



- **Feature importance**
- Ensures that predictions are chemically grounded and trustworthy
- Prediction variation with varying descriptor values
- Might encourage to consider another model



github repository



Python in the Physical Chemistry Lab

[pyPhysChem]

<https://github.com/rpoteau/pyPhysChem>



The image shows a laptop screen with several windows open, illustrating the use of Python for data analysis in a chemistry lab. The plots include a UV-vis absorption spectrum and reaction rate curves. The code snippets shown are:

```
import matplotlib.pyplot as plt
plt, data = k = matpipt('asobranc!1')
plt.plot(x, a, p, a)
```

```
import matplotlib.pyplot as plt
plt, plot = k = matpipt('asobranc!1')
plot, r, t = k = matpipt('asobranc!1')
plot, r, t = k = matpipt('asobranc!1')
```

```
k = rate_pnstant = k
plt(r = 0 = (0.215, 5, 7, 28)
plt, plt(r, plc ? 1)
```

```
l = 10 - 5
k = np.exp(-5.0)
plot, r, t = k = matpipt('asobranc!1')
plot, r, t = k = matpipt('asobranc!1')
```

Pct: Reaction Rates

plot, rate = 0, 0.0025(2, 8.0, 7, 25)